

# Generative adversarial networks and Variational AutoEncoders

**Serena Palazzo**

Data Science and Machine Learning course

November 16, 2020



THE UNIVERSITY  
*of* EDINBURGH

## Lesson from last lecture (checkpoint 5)

---

- ↪ Very good job with checkpoint 5!
- ↪ This afternoon we will do checkpoint 7
- ↪ Checkpoint 7 deadline **Friday November 20 2020**

## ↪ Generative models

- ♦ Variational auto encoders (VAE)
- ♦ Generative adversarial networks (GAN)

## ↪ Description of project 1

# Generative models

- ↪ Generative algorithms are part of unsupervised learning techniques
- ↪ A generative model can learn to mimic any distribution of it.
- ↪ They are powerful because can learn to reproduce similar models in any domain
- ↪ some of these domains are:
  - ◆ Images
  - ◆ Music
  - ◆ Speech
  - ◆ Text
  - ◆ Videos

## Discriminative vs generative algorithms

---

- ↪ **Discriminative models** aim to predict the label to which the data belongs.
  - ♦ it aims to map features to label
- ↪ **Generative models** on the other hand do the opposite: they aim to predict features given a certain label.

## Discriminative vs generative: spam emails

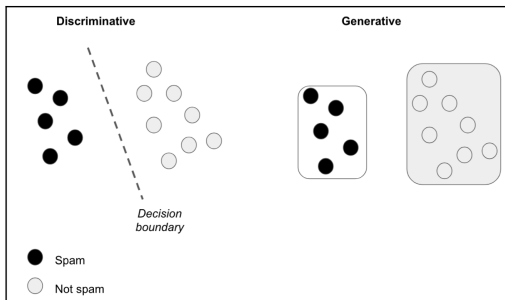
---

- ↪ Let's consider a practical example in the context of whether an email is spam or not
- ↪ We can consider  $\mathbf{x}$  the **model feature**
  - ♦ for example all the words in an email
- ↪ We can consider  $\mathbf{y}$  the **target variable**
  - ♦ whether the email is actually spam
- ↪ The discriminative and generative models will aim to answer the following questions:
  - ♦ **Discriminative  $p(\mathbf{y}|\mathbf{x})$ :** Given the input feature  $\mathbf{x}$  what is the probability of the email being spam?
  - ♦ **Generative  $p(\mathbf{x}|\mathbf{y})$ :** Given that the email is spam how likely are the input features to be  $\mathbf{x}$ ?

## Discriminative vs generative: spam emails

### In other words

- ↪ **Discriminative models** learn the boundary between classes
- ↪ **Generative models** learn to model the distribution of individual classes





## Bayes Theorem

---

- ↪ The generative model learns to predict the joint probability with the help of the following Bayes theorem:

$$p(x|y) = \frac{p(y|x)p(x)}{p(y)}$$

- ↪  $p(y|x)$  is a conditional probability, the likelihood of  $y$ , given  $x$
- ↪  $p(x|y)$  is a conditional probability, the likelihood of  $x$ , given  $y$
- ↪  $p(x)$  and  $p(y)$  are the probabilities of observing  $x$  and  $y$  independently of each other

## Examples of discriminative and generative classifiers

---

### **Discriminative classifiers**

- ↪ Logistic regression
- ↪ Nearest neighbors
- ↪ Support vector machines (SVMs)

### **Generative classifiers**

- ↪ Naive Bayes
- ↪ Hidden Markow Models (HMMs)
- ↪ Bayesian networks

# Variational Autoencoders (VAE)

## Variational autoencoders

- ↪ Variational autoencoders (VAEs) differ from the standard autoencoders.
  - ♦ they describe an observation in a **probabilistic** rather than deterministic manner.
  - ♦ The output is therefore a probability distribution rather than a single value

### Standard autoencoders vs VAEs

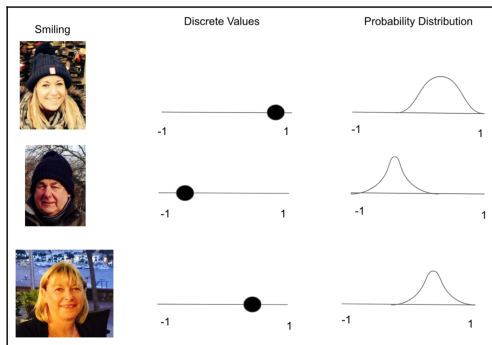
- ↪ Standard autoencoders are useful only when we want to replicate data that was input into them → limited applications in real world
- ↪ VAEs are powerful **generative** models that can be applied to cases where we don't want to output data that is the same as the input data
  - ♦ For example are useful for exploring a specific variation of input data

Glasses



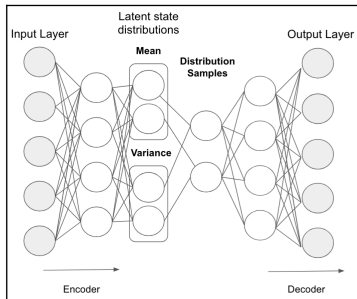
## VAE in a real-world context

- Using VAEs allows to describe each attribute in probabilistic term.
- ◆ each feature will be within a range of possible values



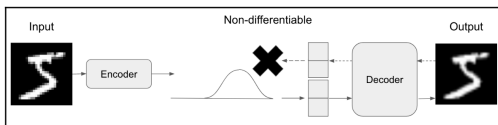
## VAE: latent feature

- ↪ The distribution of each latent attribute is sampled from the image to generate the vector that is used as input for the decoder model
- ↪ It is assumed that the distribution of each feature is Gaussian.
  - ◆ therefore 2 vectors are output: one describes the mean and the other the variance distributions



## Training VAEs

- ↪ When training a VAE it is necessary to calculate the relation of each parameter of the network wrt the overall loss, i.e. do backpropagation
- ↪ Standard Autoencoders use backpropagation, not easy for VAE since the sampling operation is not differentiable

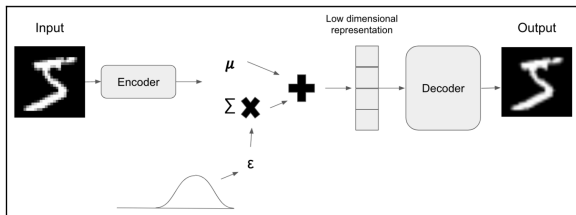


### Reparameterization trick: $z = \mu + \sigma\epsilon$

- ↪ Sample  $\epsilon$  from a unit normal distribution, shift it by mean  $\mu$  of the latent feature and scale it by the latent attributes' variance  $\sigma$

# Training VAEs

- With the introduction of the reparameterization trick, we can now train the model via simple backpropagation





## Reparameterization trick with Keras

↪ Define sampling with reparameterization trick

```
def sample_z(args):  
    mu, sigma = args  
    batch = K.shape(mu)[0]  
    dim = K.int_shape(mu)[1]  
    eps = K.random_normal(shape=(batch, dim))  
    return mu + K.exp(sigma / 2) * eps
```

↪ We then use this with a Lambda to ensure that correct gradients are computed during the backwards pass based on our values for mu and sigma

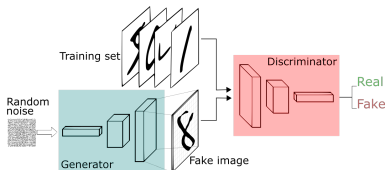
```
z = Lambda(sample_z, output_shape=(latent_dim, ), name='z')([mu,  
sigma])
```

# Generative adversarial networks

# Generative Adversarial Networks (GANs)

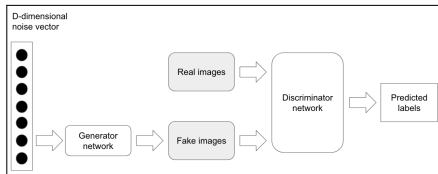
GANs are generative models that try to learn the model to generate the input distribution as realistically as possible.

- A GAN is an unsupervised learning technique and consists of two neural networks:
- ◆ A generative model **Generator** (G) generates new data points from some random uniform distribution.
  - ◆ A discriminative model **Discriminator** (D) identifies fake data produced by Generator from the real data.



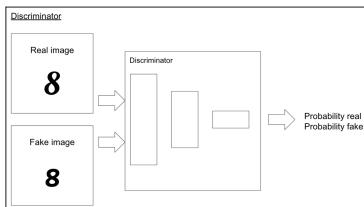
## Steps that a GAN takes

- 1 Random numbers are fed into generator and an image is generated
- 2 The generated image is fed into the discriminator along with other images taken from the real dataset
- 3 The discriminator considers all of the images fed into it and returns a probability as to whether it thinks the image is real or fake



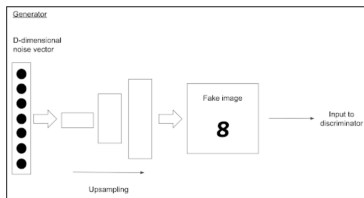
## Discriminator network

- ↪ The discriminator network is simply a standard convolutional network that categorises images being feed to it.
- ↪ It performs downsampling and classifies the images in a binary way, labeling each image as real or fake



## Generator network

- ↪ The generator network is essentially the reverse of a convolutional network.
- ↪ It takes the random noise and performs upsampling in order to output the image



## MinMax Game on GANs

- ↪ **Generator network:** try to fool the discriminator by generating real-looking images
- ↪ **Discriminator network:** try to distinguish between real and fake images

$$\min_{\theta_g} \max_{\theta_d} \left[ \mathbb{E}_{x \sim p_{data}} \log \underbrace{D_{\theta_d}(x)}_{\substack{\text{Discriminator output} \\ \text{for real data } x}} + \mathbb{E}_{z \sim p(z)} \log(1 - \underbrace{D_{\theta_d}(G_{\theta_g}(z))}_{\substack{\text{Discriminator output for} \\ \text{generated fake data } G(z)}}) \right]$$

- ↪  $D(\theta_d)$  wants to maximize objective such that  $D(x)$  is close to 1 (real) and  $D(G(z))$  is close to 0 (fake)
- ↪  $G(\theta_g)$  wants to minimize objective such that  $D(G(z))$  is close to 1
  - ♦ discriminator is fooled into thinking generated  $G(z)$  is real

## Example of Generator model in Keras

---

```
model = Sequential()  
model.add(Dense(128*7*7, input_dim = 100, activation =  
LeakyReLU(0.1)))  
model.add(BatchNormalization())  
model.add(Reshape((7,7,128)))  
model.add(UpSampling2D())  
model.add(Conv2D(64,(5,5), padding = 'same', activation =  
LeakyReLU(0.1)))  
model.add(BatchNormalization())  
model.add(UpSampling2D())  
model.add(Conv2D(1,(5,5), padding = 'same', activation = 'tanh'))
```



## Example of Discriminator model in Keras

---

```
model = Sequential()  
model.add(Conv2D(64,(5,5),padding = 'same',input_shape = (28,28,1) ,  
activation = LeakyReLU(0.1) ))  
model.add(Dropout(0.3))  
model.add(Conv2D(128,(5,5),padding = 'same', activation =  
LeakyReLU(0.1)))  
model.add(Dropout(0.3))  
model.add(Flatten())  
model.add(Dense(1,activation = 'sigmoid'))
```

Because a GAN contains two separately trained networks, its training algorithm must address 2 complications:

- ↪ GANs must juggle two different kinds of training (generator and discriminator)
- ↪ GAN convergence is hard to identify

GAN training proceeds in alternating periods:

- 1 The discriminator trains for one or more epochs
- 2 The generator trains for one or more epochs
- 3 Repeat the 2 steps to continue to train the generator and discriminator networks

## Training GANs

- ↪ We keep the generator constant during the discriminator training phase. This because the discriminator has to learn how to recognize the generator's flaws.
- ↪ We keep the discriminator constant during the generator training phase. Otherwise the generator would be trying to hit a moving target and might never converge.

### Convergence

- ↪ As the generator improves with training, the discriminator performance get worse because it cannot easily tell the difference between real and fake
- ↪ If the generator succeeds perfectly, then the discriminator has a 50% accuracy
- ↪ This poses a problem for convergence of the GAN: the discriminator feedback gets less meaningful over time
- ↪ If the GAN continues training past the point where the discriminator is giving completely random feedback, then the generator starts to train on junk feedbacks and its own quality may **collapse**

# Project 1

## Project 1: Overview and Submission

---

- ↪ An input dataset will be given and you have to perform some studies
- ↪ Perform 3-4 studies on different approaches to classify particles.
- ↪ Describe your work in the Jupyter notebook in something like a report form
- ↪ When marking the TAs will look for: (rough weights)
  - ◆ Data exploration and preprocessing (10%)
  - ◆ Model selection (30%)
  - ◆ Performance evaluation (20%)
  - ◆ Discussion, style throughout (40%)
  - ◆ Report notebooks should be submitted by **Friday December 4**

## Project 1: Standard Model

---

- ↪ A brief particle physics background useful for the dataset that you will explore for project 1

### The Standard Model

- ↪ The building blocks of matter are elementary particles.
- ↪ These particles are divided in two major types: quarks and leptons.
- ↪ The Standard Model also studies the interaction of these particles through fundamental forces (strong, weak and electromagnetic).
- ↪ For every type of particle there also exists a corresponding antiparticle.

## Project 1: elementary particles

---

- ↪ **Quarks** are fundamental constituents of matter because they combine to form hadrons. There are six quarks paired in three groups: "up/down", "charm/strange" and "top/bottom". They are held together through strong forces.
- ↪ **Hadrons** They divide in Baryons and Mesons:
  - ♦ **Baryons** are made of three quarks. For example protons are made of (uud) quarks and neutrons are made of (udd) quarks.
  - ♦ **Mesons** contain one quark and one antiquark.
- ↪ **Leptons** have a  $1/2$  spin and do not undergo strong interactions. There are six leptons, three of which have an electrical charge. These are: electron, muon and tau. The three remaining are neutrinos. A positron is the antiparticle counterpart of an electron. It possess the same mass and spin but positive charge.

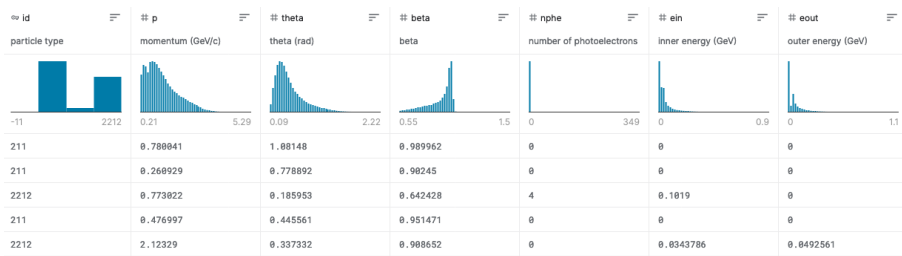
### Inelastic scattering

- ↪ Is a process used to probe the inside structure of hadrons, in this case protons.
- ↪ An incident particle (photoelectron) collides with a target proton. The kinetic energy of the incident particle is not conserved after the collision.
- ↪ During inelastic scattering a proton can break up into its constituent quarks which then form a hadronic jet. The angles of the deflection gives information about the nature of the process.



## Project 1: Dataset

- ↪ The dataset is a simulation of electron-proton inelastic scattering measured by a particle detector system.
- ↪ In order to analyse these data is necessary to recall some concepts.

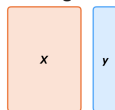


- ↪ Four particle types: positron (-11), pion (211), kaon (321), and proton (2212);
- ↪ six detector responses. Some detector responses are zero due to detector inefficiencies or incomplete geometric coverage

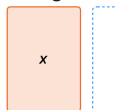
## Project 1: Training and test data

→ You will have two presplit datasets

Training data



Testing data



- Training data has truth labels ( $y$ ), testing data does not
- The problem is **multiclassification**
- You will have to classify particles type. Your **label**  $y$  will be the particle ID

## Project 1: Studies that can be performed

---

- ↪ You can perform studies based on all the techniques that you learnt during the course.
- ↪ You are free to perform studies that may or may not work. It is important to report and document everything you try and w

## Conclusions

---

- ↪ Introduced Generative models
  - ♦ Variational autoencoders
  - ♦ Generative adversarial networks
- ↪ Introduction to project 1

**Thanks for your attention and enjoy checkpoint 7 this afternoon!**